

Lossy Trapdoor Functions

by Chris Peikert, Brent Waters

by Giacomo Fenzi (supervised by Akin Ünal)

on 22 April 2021

» Motivation

Lossy Trapdoor Functions and Their Applications

Chris Peikert*
SRI International


Brent Waters†
SRI International

- * Trapdoor Functions are a basic primitive, but hard to instantiate
- * IND-CCA Security for PKE from factoring and discrete log but not lattices

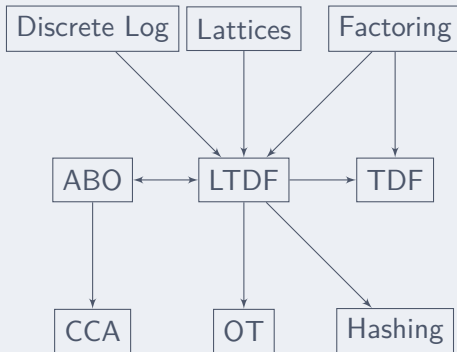
» Paper Results

- * Introduce Lossy Trapdoor Functions (LTDFs)
- * Realize LTDFs from factoring, discrete log *and* lattices
- * Show LTDFs imply TDFs
- * Black box construction of CCA-secure (witness recovering) cryptosystems, collision-resistant hash functions and oblivious transfer protocols.

» Connections



» Connections



» Trapdoor Functions

Informally, a trapdoor function is family of functions that are hard to invert without access to some additional information called a trapdoor

Definition

A trapdoor function consists of three PPT algorithms (S, F, F^{-1}) such that:

- * *Easy to sample and invert with trapdoor.*
 $S(1^\lambda) \rightarrow (s, t)$ such that $F(s, -)$ is an injective function on $\{0, 1\}^n$ and $F^{-1}(t, -)$ is its inverse
- * *Hard to invert without.* For any PPT inverter \mathcal{A} we have that $\mathcal{A}(1^\lambda, s, F(s, x))$ outputs x with negligible probability.

» Example of Trapdoor

RSA Encryption! In trapdoor form:

$S(1^\lambda)$	$F(s, x)$
Select primes $p, q, N := pq$	return $x^e \bmod N$
Select e s.t. $\gcd(e, \phi(N)) = 1$	$F^{-1}(t, y)$
$d := e^{-1} \bmod \phi(N)$	return $y^d \bmod N$
$s := (N, e)$	
$t := (s, d)$	
return (s, t)	

Correctness follows since $x^{ed} = x^1 = x$ and hardness to invert is almost exactly the RSA assumption.

Similar scheme from Pailler cryptosystem.

» Lossy Trapdoors

Informally, you either get an injective trapdoor or a 'lossy' function, and *cannot tell which is which*

Definition

An (n, k) -lossy trapdoor function consists of three PPT algorithms (S, F, F^{-1}) . We denote $S_{inj}(-) \triangleq S(-, 0)$ and $S_{lossy}(-) \triangleq S(-, 1)$.

- * *Outputs of S_{inj} are easy to compute and easy to invert with trapdoor.* $S_{inj}(1^\lambda) \rightarrow (s, t)$ s.t. that $F(s, -)$, $F^{-1}(t, -)$ are functionally as in the trapdoor case
- * *Outputs of S_{lossy} are easy to compute.* $S_{lossy}(1^\lambda) \rightarrow (s, \perp)$ s.t. $F(s, -)$ is a function on $\{0, 1\}^n$ with image size at most 2^{n-k} .
- * The description output of $S_{inj}(1^\lambda)$ and $S_{lossy}(1^\lambda)$ are computationally indistinguishable.

» Subtleties

- * The definition really relates to a collection of lossy trapdoor functions.
- * $k \triangleq k(\lambda) = \text{poly}(\lambda) \leq n$ is a parameter that represents how 'lossy' the collection is.
- * We also write $r \triangleq n - k = \text{poly}(\lambda)$ as the *residual leakage*.
- * No hardness requirement on inverting outputs of S_{inj}
- * Requirements are too strict in lattices, leads to *almost-always* lossy functions.

» All-But-One TDFs

Intuition: You have a family of functions, most of them are trapdoors, one is not. It is very hard to tell them apart.

Definition

An (n, k) -ABO TDF is a triple of PPT algorithms S, F, F^{-1} such that:

- * $S(1^\lambda, b^*) \rightarrow (s, t)$ as before
- * For any $b \neq b^*$, $F(s, b, -)$, $F^{-1}(t, b, -)$ are as in the previous definition.
- * $F(s, b^*, -)$ is a lossy function as before
- * For any b, b' the first outputs of $S(1^\lambda, b)$, $S(1^\lambda, b')$ are computationally indistinguishable.

» LTDF \implies TDF

- * Completeness: Use the injective functions generated by S_{inj} .
- * Soundness: We cannot (information theoretically) invert the lossy branch, so if we could invert the injective trapdoors we could distinguish outputs of S_{inj}, S_{lossy} , contradicting LTDF.
- * Formally, let \mathcal{A} be an inverter. We build \mathcal{D}

$$\begin{array}{l} \mathcal{D}^{\mathcal{A}}(s) \\ \hline x \leftarrow \$ \{0, 1\}^n \\ y = F(s, x) \\ x' = \mathcal{A}(s, y) \\ \mathbf{return} \ x = x' \end{array}$$

If $s \leftarrow S_{inj}(1^\lambda)$ then it succeeds nonnegligibly, while otherwise it will fail

» Realizations

We can realize LTDF from *any* encryption scheme that is:

- * Additively Homomorphic. This allows to encrypt matrices such as \mathbf{I}_n or $\mathbf{0}_n$ indistinguishably and to evaluate matrix vector products with an encrypted matrix.
- * Secure to Reuse Randomness, so that we can use the same randomness with different keys securely.
- * Isolated Randomness, so that it is only dependent on the input randomness and not on keys/messages.

We show next a realization from the DDH assumption, but a similar technique can also be employed with lattices (based on LWE) with some difficulties.

» DDH \implies LTDF

Consider the following variant of the ElGamal cryptosystem, for $m \in \{0, 1\}$, $r \in \mathbb{Z}_p$ as randomness.

$S(1^\lambda)$	$E_h(m; r)$
$\mathbb{G} \leftarrow \mathcal{G}(1^\lambda)$	return $(g^r, h^r g^m)$
$z \leftarrow \$\mathbb{Z}_p$	$D_h((c_1, c_2); r)$
$h := g^z$	return $\log_g \left(\frac{c_2}{c_1^z} \right)$
$pk := h$	
$sk := z$	
return (pk, sk)	

This scheme is semantically secure and it is additively¹ homomorphic i.e.

$$E_h(m; r) \odot E_h(m'; r') = E_h(m + m', r + r')$$

$$E_h(m; r)^x = E_h(mx; rx)$$

¹All operations done component wise

» DDH \Rightarrow LTDF

We show how to use the previous scheme to encrypt a matrix $\mathbf{M} = (m_{i,j}) \in \mathbb{Z}_p^{n \times n}$. Select n pk/sk pairs $h_i = g^{z_i}$ and n pieces of randomness r_i . Then the encryption is the matrix

$\mathbf{C} = (c_{i,j}) = (E_{h_j}(m_{i,j}; r_i))$ with the z_j s as decryption keys.

We can represent \mathbf{C} as the following two matrices:

$$\mathbf{C}_1 = \begin{bmatrix} g^{r_1} \\ \vdots \\ g^{r_n} \end{bmatrix}, \quad \mathbf{C}_2 = \begin{bmatrix} h_1^{r_1} g^{m_{1,1}} & \dots & h_n^{r_1} g^{m_{1,n}} \\ \vdots & \ddots & \vdots \\ h_1^{r_n} g^{m_{n,1}} & \dots & h_n^{r_n} g^{m_{n,n}} \end{bmatrix}$$

Via n^2 hybrid games we can show that this encryption produces indistinguishable ciphertext under DDH. We denote this operation as $\text{ME}(\mathbf{z}, \mathbf{M})$ for \mathbf{z} the vector of private keys.

» DDH \Rightarrow LTDF

We build a LTDF from the previous scheme as it follows. Note that in the injective case we encrypt the identity matrix, while in the lossy case the all zero matrix.

$S_{inj}(1^\lambda)$	$S_{lossy}(1^\lambda)$	$F_{ltdf}(\mathbf{C}, \mathbf{x})$
$\mathbb{G} \leftarrow \mathcal{G}(1^\lambda)$	$\mathbb{G} \leftarrow \mathcal{G}(1^\lambda)$	return $\mathbf{y} = \mathbf{x} \cdot \mathbf{C}$
$\mathbf{z} \leftarrow \$\mathbb{Z}_p^n$	$\mathbf{z} \leftarrow \$\mathbb{Z}_p^n$	
$\mathbf{C} \leftarrow \text{ME}(\mathbf{z}, \mathbf{I}_n)$	$\mathbf{C} \leftarrow \text{ME}(\mathbf{z}, \mathbf{0}_n)$	$F_{ltdf}^{-1}(\mathbf{z}, \mathbf{y})$
return (\mathbf{C}, \mathbf{z})	return (\mathbf{C}, \perp)	$x_i = D_{z_i}(y_i)$
		return \mathbf{x}

» DDH \Rightarrow LTDF

Subleties

- * \mathcal{G} is the group generation algorithms, it returns (G, p, g) where G is a cyclic group of prime order p with generator g . We assume DDH hardness w.r.t. \mathcal{G} .
- * \mathbf{xC} is computed by the homomorphic property. In fact, if $\mathbf{C} = \text{ME}(\mathbf{z}, \mathbf{M})$ with randomness \mathbf{r} and $h_j = g^{z_j}$

$$y_j = \bigodot_{i=1}^n c_{i,j}^{x_i} = E_{h_j}((\mathbf{xM})_j; R)$$

for R randomness that depends only on \mathbf{r} and \mathbf{x} .

- * Note that if $\mathbf{M} = \mathbf{I}_n$ then $y_j = E_{h_j}(x_j; R)$
- * If instead $\mathbf{M} = \mathbf{0}_n$ then $y_j = E_{h_j}(0; R)$

» DDH \implies LTDF

Final Checks

Now, we just have to check that the LTDF conditions are satisfied. In particular, the above construction is $(n, n - \lg p)$ -lossy.

- * The three algorithms are clearly PPT
- * A quick thought shows that the injective conditions are met
- * Indistinguishability follows from the indistinguishability of ME.
- * Finally, for outputs generated by S_{lossy} we have that $y_i = E_{h_i}(0; R)$ for some $R \in \mathbb{Z}_p$ that depends on x . R can take at most p values, the residual leakage is at most $\lg p$ and so the loss is $k = n - r \geq n - \lg p$

» LTDF \implies CCA

Requirements

We will require some primitives². We note that our cryptosystem will have message space $\{0, 1\}^\ell$.

- * We have $\Sigma = (\text{Gen}, \text{Sign}, \text{Vfy})$ a strongly unforgeable one-time signature scheme. We require that signatures are in $\{0, 1\}^v$.
- * $F = (S_{ltdf}, F_{ltdf}, F_{ltdf}^{-1})$ is an (n, k) -lossy trapdoor function.
- * $G = (S_{abo}, F_{abo}, F_{abo}^{-1})$ is an (n, k') -ABO trapdoor function with branch space $\{0, 1\}^v$.
- * \mathcal{H} is a collection of pairwise independent hash functions $\{0, 1\}^n \rightarrow \{0, 1\}^\ell$.

²All of these reduce to LTDFs

» LTDF \implies CCA

Encryption Scheme

$\mathcal{G}(1^\lambda)$	$\mathcal{E}(pk, m)$	$\mathcal{D}(sk, c)$
$(s, t) \leftarrow S_{inj}(1^\lambda)$	$(vk, sk_\sigma) = \text{Gen}(1^\lambda)$	if $\neg \text{Vfy}(vk, (c_i)_{i=1}^3, \sigma)$
$(s', t') \leftarrow S_{abo}(1^\lambda, 0^v)$	$x \leftarrow \$\{0, 1\}^n$	return \perp
$h \leftarrow \$\mathcal{H}$	$c_1 = F_{ltdf}(s, x)$	fi
$pk := (s, s', h)$	$c_2 = G_{abo}(s', vk, x)$	$x = F^{-1}(t, c_1)$
$sk := (t, t', pk)$	$c_3 = m \oplus h(x)$	if $c_1 \neq F_{ltdf}(s, x) \vee$
return (pk, sk)	$\sigma \leftarrow \text{Sign}(sk_\sigma, (c_i)_{i=1}^3)$	$c_2 \neq G_{abo}(s', vk, x)$
	return $(vk, c_1, c_2, c_3, \sigma)$	return \perp
		fi
		return $c_3 \oplus h(x)$

» LTDF \implies CCA

CCA Game

Setup is to be called once at the beginning of the game, and the attacker is allowed a single query to EncO and oracle access to DecO. The attacker wins if it outputs $b' = b$.

Setup(λ)	EncO(m_0, m_1)	DecO(c^*)
$b \leftarrow \$\{0, 1\}$	$c \leftarrow \mathcal{E}(pk, m_b)$	if $c^* \in \mathcal{T}_{enc}$
$\mathcal{T}_{enc} = \emptyset$	$\mathcal{T}_{enc} := \mathcal{T}_{enc} \cup \{c\}$	return \perp
$pk, sk \leftarrow \mathcal{G}(\lambda)$	return c	fi
return pk		return $\mathcal{D}(sk, c^*)$

» LTDF \implies CCA

Game Hops

We proceed by a sequence of games. We note that, since a single query is made to EncO we move the signature scheme generation in Setup and denote that verification key as vk^* .

$G_1(\lambda)$: This is the original CCA Security Game

$G_2(\lambda)$: In DecO if $vk = vk^*$ return \perp

$G_3(\lambda)$: In Setup choose the lossy branch of G to be vk^*

$G_4(\lambda)$: In DecO find x using G 's trapdoor rather than F 's

$G_5(\lambda)$: In Setup replace S_{inj} with S_{lossy}

The hops are as follows:

$$G_1 \approx_{\Sigma} G_2 \approx_{abo} G_3 \equiv G_4 \approx_{ltdf} G_5$$

Finally, an argument as in the TDF case shows that in G_5 even an unbounded attacker has only negligible success probability.

» Things which I did not have time to show

- * $ABO \equiv LTDF$ (see extra)
- * More efficient ABO construction from DDH
- * LTDFs from LWE
- * CPA from LTDFs
- * SUF one time signatures from LTDFs
- * UOWHFs, CRHFs from LTDFs
- * OT from LTDFs

» Related Work

- * More Constructions of LTDFs (Freeman et al.)
- * Lossy Encryption (Bellare et al.)
- * All-But-N LTDFs (Hemenway et al.)
- * All-But-Many LTDFs (Hofheinz)
- * Identity Based LTDFs (Bellare et al.)
- * Deterministic PKE (Boldyreva et al.)

Thank You!

» Notation and Entropy

- * λ is the security parameter, and we will abbreviate $n(\lambda) = \text{poly}(\lambda)$ as simply n
- * $f(-)$ denotes the function taking $x \mapsto f(x)$
- * Write $H_\infty(X)$ for the min-entropy of X . This corresponds to the optimal probability of guessing X .
- * We let $\tilde{H}_\infty(X|Y)$ be the average min-entropy of X conditioned on Y . This corresponds to the optimal probability of guessing X knowing Y .
- * We use the following lemma, if Y takes at most 2^r values then:

$$\tilde{H}_\infty(X|Y) \geq H_\infty(X) - r$$

» LTDF \implies TDF

Note that if s is generated by S_{inj} then with some non negligible probability we have that \mathcal{A} succeeds and \mathcal{D} succeeds whenever \mathcal{A} does.

Instead, if s is generated by S_{lossy} even an unbounded adversary would have best possible probability given by $2^{-\tilde{H}_\infty(x|s, F(s, x))}$. But note that $F(s, -)$ takes at most 2^r values and so by the previous lemma $\tilde{H}_\infty(x|s, F(s, x)) \geq H_\infty(x|s) - r = n - (n - k) = k$. So the probability is bounded by 2^{-k} and as such is negligible.

From the above it follows that \mathcal{D} will win the distinguishing game with non negligible probability.

» $ABO \equiv LTDF$

- * ABOs and LTDFs are equivalent.
- * $ABO \implies LTDF$. Take ABO on $\{0, 1\}$ and evaluate always on one of the branches, but switch lossy branch on generation.
- * $LTDF \implies ABO$. Generate an ABO on $\{0, 1\}$ by having $s = (s_0, s_1)$ where one of the two is lossy, and evaluation by using s_b
- * Finally, we can extend ABOs on $\{0, 1\}$ to ABOs on $\{0, 1\}^\ell$ at the cost of having residual leakage ℓr . The idea is, for lossy branch $b^* \in \{0, 1\}^\ell$, generate ℓ ABOs each with the i -th having lossy branch b_i^* .

» Pailler Cryptosystem

Composite Residuosity

- * $S(1^\lambda)$ generates $N = pq$ as a product of large primes, select g suitably, $s := (N, g)$, $t := (p, q)$
- * $F(s, x)$ splits $x = m_1 + Nm_2$ and returns $g^{m_1}m_2^N \bmod N^2$
- * $F^{-1}(t, c)$ decrypts using the factorization to compute Carmichael function